

From 100,000 profiles to a shortlist recruiters can trust

A deterministic candidate ranking system built to understand fit, reject traps, and survive reproduction.

- CPU ONLY
- NO NETWORK
- EVIDENCE-GROUNDED

84.6s native CPU runtime	0 honeypots in top 100	10/10 exact-match tier in top 10
-----------------------------	---------------------------	-------------------------------------


V7

CALIBRATED
RANKER

Fit
+ Availability
+ Integrity
+ Reasoning

The metric rewards precision where recruiters feel it

Top-10 quality carries half of the final score. The design therefore protects the first page aggressively instead of optimizing broad recall.



DESIGN CONSEQUENCE

Protect the top 10

- exact-match career evidence
- hard availability gates
- decisive honeypot removal
- deterministic tie-breaking

The JD is a systems-thinking problem, not a keyword query

DEPTH

Embeddings, retrieval, ranking, vector search, evaluation

Career evidence over skill-list density

SHIPPER MODE

Production ownership, product judgment, recent coding

Deployed outcomes over research-only claims

FIT

Product-company context, stable trajectory, India logistics

Role intent over title prestige


HIRABILITY

Recent activity, response behavior, interview reliability

Reachability over paper perfection

TRAP: perfect AI keywords + contradictory career or stale behavior = not a strong candidate

One auditable score, with explicit gates around it



CALIBRATED FIT WEIGHTS



The biggest gains came from removing false confidence

The initial system looked sophisticated, but several mechanics were not defensible. Each was replaced with a simpler measurable behavior.

AREA	BEFORE	AFTER	WHY IT MATTERS
Integrity	29,220 profiles penalized	57 decisive contradictions	Fixed current-role date-span false positives
Behavior	Most strong profiles saturated at max	Hundreds of distinct values	Removed duplicate counting; added reachability gates
Semantics	JD embedding included the whole bundle	Focused deterministic evidence	Default no longer depends on ignored 154 MB artifact
Combination	Three linear strategies called an ensemble	One calibrated weight vector	Mathematically honest and easier to defend

Three traps, three explicit defenses

KEYWORD STUFFER

Perfect AI skills; career says marketing or unrelated work

DEFENSE

Title-career fit + skill-career validation + stuffing penalty

DOWN-RANK

IMPOSSIBLE PROFILE

Experience, dates, or expert-skill durations contradict each other

DEFENSE

Cross-field integrity scorer; decisive contradictions only

REMOVE

UNREACHABLE STAR

Excellent retrieval background; 5% response rate and stale activity

DEFENSE

Weighted availability score + explicit unreachable gate

DOWN-WEIGHT

Result: zero detected honeypots and zero very-low-availability profiles in the submitted top 100.

The full-pool run stays fast and the shortlist stays clean

100,000

profiles ranked

84.6s

native CPU time

57


decisive contradictions

100/100

unique reasonings

NATIVE RUNTIME BREAKDOWN

Load candidates



9.0s

Score + rank



74.7s

Budget used

28%

of 5-minute limit

Docker reproduction: 162.8s · identical output SHA-256

SHORTLIST AUDIT

Top 10 exact-match tier

10 / 10

Top 50 tier 5 + tier 4

13 + 37

Top 100 tier 5 + tier 4

13 + 87

Honeypots / unavailable

0 / 0

Built to survive every evaluation stage, not just scoring

1

FORMAT

Official validator passes; exactly 100 ranked rows

2

SCORE

Top-10-first ranking with exact-match prior and evidence

3

REPRODUCE

Native: 84.6s | Docker reproduction: 162.8s | no network

4

MANUAL

100 unique, candidate-grounded reasons; honest trade-offs

5

DEFEND

One calibrated score; explicit gates; documented limitations

HONEST LIMITATION

No public labels means no one can know the hidden score before judging.

Our response is disciplined inference: challenge-specific priors, measurable audits, deterministic reproduction, and clear reasoning.

READY TO SUBMIT. READY TO EXPLAIN.